

# Unit 2.1: Understanding Parallelism I

## Video lesson 2 – Expressing tasks

Eduard Ayguadé, Josep Ramon Herrero,  
Daniel Jiménez and Gladys Utrera

Barcelona Supercomputing Center  
Universitat Politècnica de Catalunya

Let's go back to the car retail database example.


| ID#  | Model   | Year | Color | Dealer | Price    |
|------|---------|------|-------|--------|----------|
| 4523 | Civic   | 2002 | Blue  | MN     | \$18,000 |
| 3476 | Corolla | 1999 | White | IL     | \$15,000 |
| 7623 | Camry   | 2001 | Green | NY     | \$21,000 |
| 9834 | Prius   | 2001 | Green | CA     | \$18,000 |
| 6734 | Civic   | 2001 | White | OR     | \$17,000 |
| 5342 | Altima  | 2001 | Green | FL     | \$19,000 |
| 3845 | Maxima  | 2001 | Blue  | NY     | \$22,000 |
| 8354 | Accord  | 2000 | Green | VT     | \$18,000 |
| 4395 | Civic   | 2001 | Red   | CA     | \$17,000 |
| 7352 | Civic   | 2002 | Red   | WA     | \$18,000 |

And assume that we want to count how many Green cars are available to sell.

# Motivation – I

- One could traverse all the records  $X[0] \dots X[n-1]$  in the database  $X$  and check if the `Color` field matches the required value `Green`, storing the number of matches in variable `count`

database  $X$




- A possible sequential program could be:

```
count = 0;  
for ( i = 0 ; i < n ; i++ )  
    if (X[i].Color == "Green") count++;
```

whose computation time on a single processor would be proportional to the number of records in the database  $T_1 \propto n$

- One could divide the traversal in  $P$  groups (tasks), for example for  $P = 4$ :


database X



checking the Color field for a subset of  $n \div P$  consecutive records, and counting on a per-task "private" copy of variable count

- However, we still need to "globally" count the number of records found that match the condition by combining the individual "private" counts into the original count variable

database X



- Up to this point you could anticipate that the computation time would be divided by the number of tasks  $P$  if  $P$  workers are used to do the computation

$$T_P = T_1 \div P$$

with an additional "overhead" to perform this global reduction

$$T_p = T_1 \div P + T_{ovh}(P)$$

probably proportional to the number of workers  $P$

Consider now that we want to execute the following query:

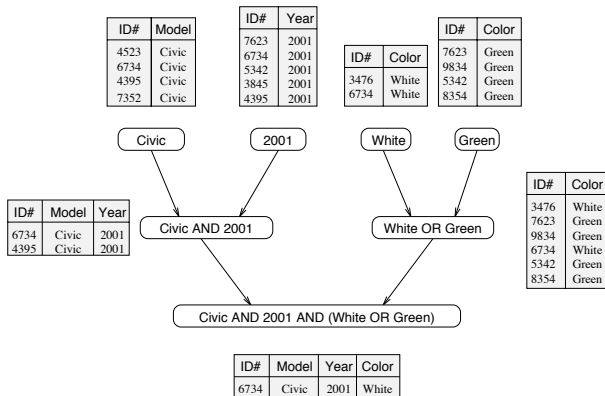
```
MODEL = 'CIVIC' AND YEAR = 2001 AND  
(COLOR = 'GREEN' OR COLOR = 'WHITE')
```

on our car dealer database:

| ID#  | Model   | Year | Color | Dealer | Price    |
|------|---------|------|-------|--------|----------|
| 4523 | Civic   | 2002 | Blue  | MN     | \$18,000 |
| 3476 | Corolla | 1999 | White | IL     | \$15,000 |
| 7623 | Camry   | 2001 | Green | NY     | \$21,000 |
| 9834 | Prius   | 2001 | Green | CA     | \$18,000 |
| 6734 | Civic   | 2001 | White | OR     | \$17,000 |
| 5342 | Altima  | 2001 | Green | FL     | \$19,000 |
| 3845 | Maxima  | 2001 | Blue  | NY     | \$22,000 |
| 8354 | Accord  | 2000 | Green | VT     | \$18,000 |
| 4395 | Civic   | 2001 | Red   | CA     | \$17,000 |
| 7352 | Civic   | 2002 | Red   | WA     | \$18,000 |




- A possible query plan could be:



- Each of these operations in the query plan could be a task, each computing an intermediate table of entries that satisfy particular conditions
- Are they independent?
  - Some of them are, for example tasks "*Civic*", "*2001*", "*Green*" and "*White*"
  - Others are not independent. For example, task "*Civic\_AND\_2001*" can not start its execution until both tasks "*Civic*" and "*2001*" complete

# Tasks and dependences

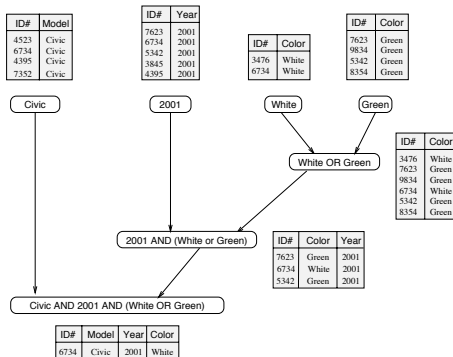
- Dependences impose task execution ordering constraints that need to be fulfilled in order to guarantee correct results




- **Task dependence graph:** graphical representation of the task decomposition

# Tasks and dependences

- Other query plans are possible, for example




... with different task dependence graphs and potential to execute tasks in parallel



- **Task dependence graph** abstraction

- Directed Acyclic Graph
- Node = task, its weight represents the amount of work to be done
- Edge = dependence, i.e. successor node can only execute after predecessor node has completed

# Computing the parallelism




- Parallel machine abstraction

- P identical processors
- Each processor executes a node at a time


- $$T_1 = \sum_{i=1}^{nodes} (work\_node_i)$$

# Computing the parallelism



- Critical path: path in the task graph with the highest accumulated work
- $T_{\infty} = \sum_{i \in \text{criticalpath}} (\text{work\_node}_i)$ , assuming sufficient processors


# Computing the parallelism



- $Parallelism = T_1/T_\infty$ , if sufficient processors were available
- $P_{min}$  is the minimum number of processors necessary to achieve *Parallelism*



# Computing the parallelism



- $T_1 = 47$  including tasks {ABCDEFGH}

- Possible paths:

{ABCEH} with total cost 29  
{ABDFH} with total cost 40  
{ABDGH} with total cost 33



- $T_\infty = 40$  for critical path {ABDFH}

- $Parallelism = 47/40 = 1.175$

- $P_{min} = 2$

# Going back to our database query example

Consider the task dependency graphs for the two database queries, assuming *work\_node* is proportional to the number of inputs to be processed



Which are  $T_1$ ,  $T_\infty$  and *Parallelism* in each case?

# Going back to our database query example

- For graph (a) the critical path is determined by any of the two task sequences  $\{3, 6, 7\}$  or  $\{4, 6, 7\}$

$$T_1^{(a)} = 63, T_\infty^{(a)} = 27, \text{Parallelism}^{(a)} = 63/27 = 2.33$$

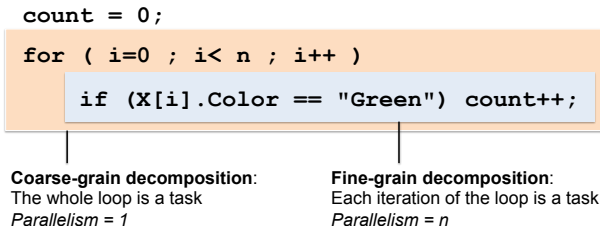
- However, for graph (b) the critical path is determined by any of the two task sequences  $\{1, 5, 6, 7\}$  or  $\{2, 5, 6, 7\}$

$$T_1^{(b)} = 64, T_\infty^{(b)} = 34, \text{Parallelism}^{(b)} = 64/34 = 1.88$$

- Do both query plans require the same minimum number of processors  $P_{min}$  to achieve this potential *Parallelism*?

# Granularity and parallelism


- The granularity of the task decomposition is determined by the computational size of the nodes (tasks) in the task graph
- Example: counting matches in our car dealer database



# Granularity and parallelism

- Example: A task could be in charge of checking a number of consecutive elements  $m$  of the database:

database X



with a potential *parallelism*  $= n \div m$

$$1 < (n \div m) < n$$

- It would appear that the parallelism is higher when going to fine-grain task decompositions

|                     | Coarse grain | Fine grain | Medium grain |
|---------------------|--------------|------------|--------------|
| Number of tasks     | 1            | n          | n / m        |
| Iterations per task | n            | 1          | m            |
| Parallelism         | 1            | n          | n / m        |

- However, tradeoff between potential parallelism and overheads related with its exploitation (e.g. creation of tasks, synchronization, exchange of data, ...)

# Unit 2.1: Understanding Parallelism I

## Video lesson 2 – Expressing tasks

Eduard Ayguadé, Josep Ramon Herrero,  
Daniel Jiménez and Gladys Utrera

Barcelona Supercomputing Center  
Universitat Politècnica de Catalunya